

# **Module 3**

## **Memory management fundamentals**

# Memory management

- we can manage memory **manually** or **automatically** 😊

# Memory management

- we can manage memory **manually** or **automatically** 😊
- if manually:

# Memory management

- we can manage memory **manually** or **automatically** 😊
- if manually:
  - we create objects but we need also to delete them

# Memory management

- we can manage memory **manually** or **automatically** 😊
- if manually:
  - we create objects but we need also to delete them
  - typically, creation is supported by some **allocator** - it's *"give me a new memory"* rather than *"give me **this** memory region"*

# Memory management

- we can manage memory **manually** or **automatically** 😊
- if manually:
  - we create objects but we need also to delete them
  - typically, creation is supported by some **allocator** - it's *"give me a new memory"* rather than *"give me **this** memory region"*
- if automatically:

# Memory management

- we can manage memory **manually** or **automatically** 😊
- if manually:
  - we create objects but we need also to delete them
  - typically, creation is supported by some **allocator** - it's *"give me a new memory"* rather than *"give me **this** memory region"*
- if automatically:
  - we only create objects and don't need to delete them - they will be automatically deleted

# Memory management

- we can manage memory **manually** or **automatically** 😊
- if manually:
  - we create objects but we need also to delete them
  - typically, creation is supported by some **allocator** - it's *"give me a new memory"* rather than *"give me **this** memory region"*
- if automatically:
  - we only create objects and don't need to delete them - they will be automatically deleted
  - may be handled at:
    - **execution** time - both creation and reclamation are supported by **allocator** and some **garbage collection** mechanism
    - **compilation time** - discovering what should be reclaimed is discovered during compilation time



# Memory management

- we can manage memory **manually** or **automatically** 😊
- if manually:
  - we create objects but we need also to delete them
  - typically, creation is supported by some **allocator** - it's *"give me a new memory"* rather than *"give me **this** memory region"*
- if automatically:
  - we only create objects and don't need to delete them - they will be automatically deleted
  - may be handled at:
    - **execution** time - both creation and reclamation are supported by **allocator** and some **garbage collection** mechanism
    - **compilation time** - discovering what should be reclaimed is discovered during compilation time (which is not trivial, we will see it soon!)

# Manual memory management

Also known as *explicit allocation/deallocation*:

```
#include<stdio.h>
int main()
{
    int *ptr;
    ptr = (int*)malloc(sizeof(int));
    if (ptr == 0)
    {
        printf("ERROR: Out of memory\n");
        return 1;
    }
    *ptr = 25;
    printf("%d\n", *ptr);
    free(ptr);
    return 0;
}
```

Problems:

# Manual memory management

Also known as *explicit allocation/deallocation*:

```
#include<stdio.h>
int main()
{
    int *ptr;
    ptr = (int*)malloc(sizeof(int));
    if (ptr == 0)
    {
        printf("ERROR: Out of memory\n");
        return 1;
    }
    *ptr = 25;
    printf("%d\n", *ptr);
    free(ptr);
    return 0;
}
```

Problems:

- tiresome - a lot of ceremony code, instead of business code

# Manual memory management

Also known as *explicit allocation/deallocation*:

```
#include<stdio.h>
int main()
{
    int *ptr;
    ptr = (int*)malloc(sizeof(int));
    if (ptr == 0)
    {
        printf("ERROR: Out of memory\n");
        return 1;
    }
    *ptr = 25;
    printf("%d\n", *ptr);
    free(ptr);
    return 0;
}
```

Problems:

- tiresome - a lot of ceremony code, instead of business code
- memory leak - if we forget to call **free** (🙄)

# Manual memory management

Also known as *explicit allocation/deallocation*:

```
#include<stdio.h>
int main()
{
    int *ptr;
    ptr = (int*)malloc(sizeof(int));
    if (ptr == 0)
    {
        printf("ERROR: Out of memory\n");
        return 1;
    }
    *ptr = 25;
    printf("%d\n", *ptr);
    free(ptr);
    return 0;
}
```

Problems:

- tiresome - a lot of ceremony code, instead of business code
- memory leak - if we forget to call **free** (🙄)
- dangling pointer - if we use a pointer after calling **free** (🙄🙄)

# Manual memory management

Also known as *explicit allocation/deallocation*:

```
#include<stdio.h>
int main()
{
    int *ptr;
    ptr = (int*)malloc(sizeof(int));
    if (ptr == 0)
    {
        printf("ERROR: Out of memory\n");
        return 1;
    }
    *ptr = 25;
    printf("%d\n", *ptr);
    free(ptr);
    return 0;
}
```

Problems:

- tiresome - a lot of ceremony code, instead of business code
- memory leak - if we forget to call **free** (🙄)
- dangling pointer - if we use a pointer after calling **free** (🙄🙄)
- cross-library support - every library promotes its own "helpers" to handle that (aka *smart pointers*)

# Automatic memory management

Addresses common manual memory management problems:

# Automatic memory management

Addresses common manual memory management problems:

- tiresome - only business code, no ceremony



# Automatic memory management

Addresses common manual memory management problems:

- tiresome - only business code, no ceremony
- memory leak - should not happen, "**free** calls" are somehow automagically injected

# Automatic memory management

Addresses common manual memory management problems:

- tiresome - only business code, no ceremony
- memory leak - should not happen, "**free** calls" are somehow automagically injected
- dangling pointer - pointers/references are managed and it should not happen

# Automatic memory management

Addresses common manual memory management problems:

- tiresome - only business code, no ceremony
- memory leak - should not happen, "**free** calls" are somehow automagically injected
- dangling pointer - pointers/references are managed and it should not happen
  - in other words, regular code should not see a reference **after** the memory has been reclaimed

# Automatic memory management

Addresses common manual memory management problems:

- tiresome - only business code, no ceremony
- memory leak - should not happen, "**free** calls" are somehow automagically injected
- dangling pointer - pointers/references are managed and it should not happen
  - in other words, regular code should not see a reference **after** the memory has been reclaimed
- cross-library support - *lingua franca* of the whole ecosystem

# Automatic memory management

Addresses common manual memory management problems:

- tiresome - only business code, no ceremony
- memory leak - should not happen, "**free** calls" are somehow automagically injected
- dangling pointer - pointers/references are managed and it should not happen
  - in other words, regular code should not see a reference **after** the memory has been reclaimed
- cross-library support - *lingua franca* of the whole ecosystem

But... adds some new problems:

# Automatic memory management

Addresses common manual memory management problems:

- tiresome - only business code, no ceremony
- memory leak - should not happen, "**free** calls" are somehow automagically injected
- dangling pointer - pointers/references are managed and it should not happen
  - in other words, regular code should not see a reference **after** the memory has been reclaimed
- cross-library support - *lingua franca* of the whole ecosystem

But... adds some new problems:

- less or more overhead - "automagic" have some costs:

# Automatic memory management

Addresses common manual memory management problems:

- tiresome - only business code, no ceremony
- memory leak - should not happen, "**free** calls" are somehow automagically injected
- dangling pointer - pointers/references are managed and it should not happen
  - in other words, regular code should not see a reference **after** the memory has been reclaimed
- cross-library support - *lingua franca* of the whole ecosystem

But... adds some new problems:

- less or more overhead - "automagic" have some costs:
  - in case of *execution time* - we need some **runtime/garbage collection** process, and it will just consume resources and influence (pause? slowdown?) the program

# Automatic memory management

Addresses common manual memory management problems:

- tiresome - only business code, no ceremony
- memory leak - should not happen, "**free calls**" are somehow automagically injected
- dangling pointer - pointers/references are managed and it should not happen
  - in other words, regular code should not see a reference **after** the memory has been reclaimed
- cross-library support - *lingua franca* of the whole ecosystem

But... adds some new problems:

- less or more overhead - "automagic" have some costs:
  - in case of *execution time* - we need some **runtime/garbage collection** process, and it will just consume resources and influence (pause? slowdown?) the program
  - in case of *compile time* - smart enough compiler will be just slower (and probably still need some *ceremony code* from the developer)



# Automatic memory management

Addresses common manual memory management problems:

- tiresome - only business code, no ceremony
- memory leak - should not happen, "**free calls**" are somehow automagically injected
- dangling pointer - pointers/references are managed and it should not happen
  - in other words, regular code should not see a reference **after** the memory has been reclaimed
- cross-library support - *lingua franca* of the whole ecosystem

But... adds some new problems:

- less or more overhead - "automagic" have some costs:
  - in case of *execution time* - we need some **runtime/garbage collection** process, and it will just consume resources and influence (pause? slowdown?) the program
  - in case of *compile time* - smart enough compiler will be just slower (and probably still need some *ceremony code* from the developer)
- it is an abstraction of infinite memory:

# Automatic memory management

Addresses common manual memory management problems:

- tiresome - only business code, no ceremony
- memory leak - should not happen, "**free calls**" are somehow automagically injected
- dangling pointer - pointers/references are managed and it should not happen
  - in other words, regular code should not see a reference **after** the memory has been reclaimed
- cross-library support - *lingua franca* of the whole ecosystem

But... adds some new problems:

- less or more overhead - "automagic" have some costs:
  - in case of *execution time* - we need some **runtime/garbage collection** process, and it will just consume resources and influence (pause? slowdown?) the program
  - in case of *compile time* - smart enough compiler will be just slower (and probably still need some *ceremony code* from the developer)
- it is an abstraction of infinite memory:
  - sometimes it "leaks" - implementation details (sometimes complex and unclear) may become problematic in edge cases and non trivial to analyze

# Automatic memory management

Addresses common manual memory management problems:

- tiresome - only business code, no ceremony
- memory leak - should not happen, "**free** calls" are somehow automagically injected
- dangling pointer - pointers/references are managed and it should not happen
  - in other words, regular code should not see a reference **after** the memory has been reclaimed
- cross-library support - *lingua franca* of the whole ecosystem

But... adds some new problems:

- less or more overhead - "automagic" have some costs:
  - in case of *execution time* - we need some **runtime/garbage collection** process, and it will just consume resources and influence (pause? slowdown?) the program
  - in case of *compile time* - smart enough compiler will be just slower (and probably still need some *ceremony code* from the developer)
- it is an abstraction of infinite memory:
  - sometimes it "leaks" - implementation details (sometimes complex and unclear) may become problematic in edge cases and non trivial to analyze
  - it is impossible to produce solution good enough for all possible scenarios - we need to start thinking how to choose/configure it

# Automatic memory management

Addresses common manual memory management problems:

- tiresome - only business code, no ceremony
- memory leak - should not happen, "**free calls**" are somehow automagically injected
- dangling pointer - pointers/references are managed and it should not happen
  - in other words, regular code should not see a reference **after** the memory has been reclaimed
- cross-library support - *lingua franca* of the whole ecosystem

But... adds some new problems:

- less or more overhead - "automagic" have some costs:
  - in case of *execution time* - we need some **runtime/garbage collection** process, and it will just consume resources and influence (pause? slowdown?) the program
  - in case of *compile time* - smart enough compiler will be just slower (and probably still need some *ceremony code* from the developer)
- it is an abstraction of infinite memory:
  - sometimes it "leaks" - implementation details (sometimes complex and unclear) may become problematic in edge cases and non trivial to analyze
  - it is impossible to produce solution good enough for all possible scenarios - we need to start thinking how to choose/configure it
- BTW, **that's why we are here in this course** 😊

# Automatic memory management

Two groups of algorithms:

- reference counting (and similar)
- tracing garbage collector

# Reference counting

# Reference counting

- simple idea:

# Reference counting

- simple idea:
  - automatically **counts** how many references/pointers point to an object



# Reference counting

- simple idea:
  - automatically **counts** how many references/pointers point to an object
  - when the count drops to 0 - it's high time we removed an object

# Reference counting

- simple idea:
  - automatically **counts** how many references/pointers point to an object
  - when the count drops to 0 - it's high time we removed an object
- implementations:

# Reference counting

- simple idea:
  - automatically **counts** how many references/pointers point to an object
  - when the count drops to 0 - it's high time we removed an object
- implementations:
  - *smart pointers* without the runtime - C++, COM, Rust:

```
int main()
{
    std::shared_ptr<Foo> shared(new Foo);
    std::cout << shared.use_count() << '\n';
}
```

# Reference counting

- simple idea:
  - automatically **counts** how many references/pointers point to an object
  - when the count drops to 0 - it's high time we removed an object
- implementations:
  - *smart pointers* without the runtime - C++, COM, Rust:

```
int main()
{
    std::shared_ptr<Foo> shared(new Foo);
    std::cout << shared.use_count() << '\n';
}
```

- *smart pointers* with the runtime - Python

# Reference counting

- advantages:

# Reference counting

- advantages:
  - deterministic deallocation - as soon we get 0, we can reclaim memory **immediately**

# Reference counting

- advantages:
  - deterministic deallocation - as soon we get 0, we can reclaim memory **immediately**
  - incremental work - we clean up **constantly**, no big "churns" of work

# Reference counting

- advantages:
  - deterministic deallocation - as soon we get 0, we can reclaim memory **immediately**
  - incremental work - we clean up **constantly**, no big "churns" of work
  - simple implementation - at least until not doing something more sophisticated:



# Reference counting

- advantages:
  - deterministic deallocation - as soon we get 0, we can reclaim memory **immediately**
  - incremental work - we clean up **constantly**, no big "churns" of work
  - simple implementation - at least until not doing something more sophisticated:
    - [deferred reference counting](#) - avoiding counting reference stored on the stack
    - RC Immix - [Taking Off the Gloves with Reference Counting Immix](#) paper, '2013

# Reference counting

- advantages:
  - deterministic deallocation - as soon we get 0, we can reclaim memory **immediately**
  - incremental work - we clean up **constantly**, no big "churns" of work
  - simple implementation - at least until not doing something more sophisticated:
    - [deferred reference counting](#) - avoiding counting reference stored on the stack
    - RC Immix - [Taking Off the Gloves with Reference Counting Immix](#) paper, '2013
- disadvantages:

# Reference counting

- advantages:
  - deterministic deallocation - as soon we get 0, we can reclaim memory **immediately**
  - incremental work - we clean up **constantly**, no big "churns" of work
  - simple implementation - at least until not doing something more sophisticated:
    - [deferred reference counting](#) - avoiding counting reference stored on the stack
    - RC Immix - [Taking Off the Gloves with Reference Counting Immix](#) paper, '2013
- disadvantages:
  - overhead - every assignment between references has a cost now

# Reference counting

- advantages:
  - deterministic deallocation - as soon we get 0, we can reclaim memory **immediately**
  - incremental work - we clean up **constantly**, no big "churns" of work
  - simple implementation - at least until not doing something more sophisticated:
    - [deferred reference counting](#) - avoiding counting reference stored on the stack
    - RC Immix - [Taking Off the Gloves with Reference Counting Immix](#) paper, '2013
- disadvantages:
  - overhead - every assignment between references has a cost now  
(although "recent work narrowed the gap between reference counting and the fastest tracing collectors to **within 10%**")

# Reference counting

- advantages:
  - deterministic deallocation - as soon we get 0, we can reclaim memory **immediately**
  - incremental work - we clean up **constantly**, no big "churns" of work
  - simple implementation - at least until not doing something more sophisticated:
    - [deferred reference counting](#) - avoiding counting reference stored on the stack
    - RC Immix - [Taking Off the Gloves with Reference Counting Immix](#) paper, '2013
- disadvantages:
  - overhead - every assignment between references has a cost now  
(although "recent work narrowed the gap between reference counting and the fastest tracing collectors to **within 10%**")
  - cyclic references - needs some additional help (overhead)

# Reference counting on steroids

- *affine type system* - every variable (resource) is used **at most once**, which is enforced by the compiler

# Reference counting on steroids

- *affine type system* - every variable (resource) is used **at most once**, which is enforced by the compiler
- the compiler limits the usage of resources - in a way that allows to track resources usage

# Reference counting on steroids

- *affine type system* - every variable (resource) is used **at most once**, which is enforced by the compiler
- the compiler limits the usage of resources - in a way that allows to track resources usage
  - in other words, if there is **always** at most only one owner of data, we know we don't need it when the owner dies!



# Reference counting on steroids

- *affine type system* - every variable (resource) is used **at most once**, which is enforced by the compiler
- the compiler limits the usage of resources - in a way that allows to track resources usage
  - in other words, if there is **always** at most only one owner of data, we know we don't need it when the owner dies!
- Rust is a new unmanaged language (without the runtime) with affine type system - thanks to techniques like **move semantics**, **ownership** and **borrowing**

# Affine type system - Rust

*Move semantics* example:

```
let s1 = String::from("hello");  
let s2 = s1;  
  
println!("{}", world!", s1);
```

# Affine type system - Rust

*Move semantics* example:

```
let s1 = String::from("hello");
let s2 = s1;

println!("{}", world!", s1);
```

```
error[E0382]: borrow of moved value: `s1`
--> src/main.rs:4:28
  |
2 |   let s1 = String::from("hello");
  |   -- move occurs because `s1` has type `String`, which does not implement the `Copy` trait
3 |   let s2 = s1;
  |   -- value moved here
4 |   println!("{}", world!", s1);
  |                          ^^ value borrowed here after move
```

# Affine type system - Rust

*Move semantics* example:

```
fn main() {  
    let s = String::from("hello"); // s comes into scope  
    takes_ownership(s);           // s value moves into the function...  
                                   // ... and so is no longer valid here  
  
}  
fn takes_ownership(some_string: String) { // some_string comes into scope  
    println!("{}", some_string);  
} // Here, some_string goes out of scope and `drop` is called. The backing  
   // memory is freed.
```

# Affine type system - Rust

*Borrowing* example:

```
fn main() {  
    let s = String::from("hello");  
    use_string(&s);  
    // string goes of scope here  
}  
  
fn use_string(some_string: &String) {  
    // string has been borrowed, increases it's "counter usage"  
}
```

# Affine type system - Rust

*Borrowing* example:

```
fn main() {  
    let s = String::from("hello");  
    change(&s);  
}  
  
fn change(some_string: &String) {  
    some_string.push_str(", world");  
}
```

# Affine type system - Rust

*Borrowing* example:

```
fn main() {  
    let s = String::from("hello");  
    change(&s);  
}  
  
fn change(some_string: &String) {  
    some_string.push_str(", world");  
}
```

```
error[E0596]: cannot borrow `*some_string` as mutable, as it is behind a `&` reference  
--> src/main.rs:8:5  
7 | fn use_string(some_string: &String) {  
  |                   ----- help: consider changing this to be a mutable reference: `&mut String`  
8 |     some_string.push_str(", world");  
  |     ^^^^^^^^^^^^^^^ `some_string` is a `&` reference, so the data it refers to cannot be borrowed as mutable
```

# Affine type system - Rust

*Borrowing* example:

```
fn main() {  
    let mut s = String::from("hello");  
    change(&mut s);  
}  
  
fn change(some_string: &mut String) {  
    some_string.push_str(", world");  
}
```



# Affine type system - Rust

*Borrowing* example:

```
let mut s = String::from("hello");  
  
let r1 = &mut s;  
let r2 = &mut s;  
  
println!("{}", {}, r1, r2);
```

# Affine type system - Rust

*Borrowing* example:

```
let mut s = String::from("hello");

let r1 = &mut s;
let r2 = &mut s;

println!("{}", {}, r1, r2);
```

```
error[E0499]: cannot borrow `s` as mutable more than once at a time
--> src/main.rs:4:14
3 |   let r1 = &mut s;
  |           ----- first mutable borrow occurs here
4 |   let r2 = &mut s;
  |           ^^^^^^^ second mutable borrow occurs here
5 |   println!("{}", {}, r1, r2);
  |                       -- first borrow later used here
```

# Affine type system - Rust

*Borrowing* example:

```
let mut s = String::from("hello");

let r1 = &s; // no problem
let r2 = &s; // no problem
let r3 = &mut s; // BIG PROBLEM

println!("{}", {}, and {}", r1, r2, r3);
```

# Affine type system - Rust

*Borrowing* example:

```
let mut s = String::from("hello");

let r1 = &s; // no problem
let r2 = &s; // no problem
let r3 = &mut s; // BIG PROBLEM

println!("{}", {}, and {}", r1, r2, r3);
```

```
error[E0502]: cannot borrow `s` as mutable because it is also borrowed as immutable
--> src/main.rs:5:14
```

```
3 |     let r1 = &s; // no problem
  |               -- immutable borrow occurs here
4 |     let r2 = &s; // no problem
5 |     let r3 = &mut s; // BIG PROBLEM
  |               ^^^^^^ mutable borrow occurs here
6 |     println!("{}", {}, and {}", r1, r2, r3);
  |                                   -- immutable borrow later used here
```

# Tracing garbage collector

# Tracing garbage collector

- "simple" idea - **trace** which objects are used, and **collect** (delete) those which are not (☹)

# Tracing garbage collector

- "simple" idea - **trace** which objects are used, and **collect** (delete) those which are not (☹️)
- many aspects of implementation:

# Tracing garbage collector

- "simple" idea - **trace** which objects are used, and **collect** (delete) those which are not (☹)
- many aspects of implementation:
  - how & when to trace?



# Tracing garbage collector

- "simple" idea - **trace** which objects are used, and **collect** (delete) those which are not (☹)
- many aspects of implementation:
  - how & when to trace?
  - what does "*object is used*" mean?

# Tracing garbage collector

- "simple" idea - **trace** which objects are used, and **collect** (delete) those which are not (☹)
- many aspects of implementation:
  - how & when to trace?
  - what does "*object is used*" mean?
  - how & when to reclaim?

# Tracing garbage collector

- "simple" idea - **trace** which objects are used, and **collect** (delete) those which are not (☹)
- many aspects of implementation:
  - how & when to trace?
  - what does "*object is used*" mean?
  - how & when to reclaim?
- we have many tracing GCs implementation in managed runtimes (.NET, JVM, Ruby, Go, ...)

# Tracing garbage collector

- "simple" idea - **trace** which objects are used, and **collect** (delete) those which are not (☹)
- many aspects of implementation:
  - how & when to trace?
  - what does "*object is used*" mean?
  - how & when to reclaim?
- we have many tracing GCs implementation in managed runtimes (.NET, JVM, Ruby, Go, ...)
- the concept is old:

# Tracing garbage collector

- "simple" idea - **trace** which objects are used, and **collect** (delete) those which are not (☹️)
- many aspects of implementation:
  - how & when to trace?
  - what does *"object is used"* mean?
  - how & when to reclaim?
- we have many tracing GCs implementation in managed runtimes (.NET, JVM, Ruby, Go, ...)
- the concept is old:
  - John McCarthy. *Recursive functions of symbolic expressions and their computation by machine.*  
- 1958

# Tracing garbage collector

- "simple" idea - **trace** which objects are used, and **collect** (delete) those which are not (☹)
- many aspects of implementation:
  - how & when to trace?
  - what does "object is used" mean?
  - how & when to reclaim?
- we have many tracing GCs implementation in managed runtimes (.NET, JVM, Ruby, Go, ...)
- the concept is old:
  - John McCarthy. *Recursive functions of symbolic expressions and their computation by machine.* - 1958
  - "**John McCarthy** - American computer scientist and cognitive scientist. McCarthy was **one of the founders of the discipline of artificial intelligence.** (...), **developed the Lisp programming language** (...) **significantly influenced the design of the ALGOL programming language,** (...) and **invented garbage collection.**"

# Tracing garbage collector

- "simple" idea - **trace** which objects are used, and **collect** (delete) those which are not (☹)
- many aspects of implementation:
  - how & when to trace?
  - what does "object is used" mean?
  - how & when to reclaim?
- we have many tracing GCs implementation in managed runtimes (.NET, JVM, Ruby, Go, ...)
- the concept is old:
  - John McCarthy. *Recursive functions of symbolic expressions and their computation by machine.* - 1958
  - "**John McCarthy** - American computer scientist and cognitive scientist. McCarthy was **one of the founders of the discipline of artificial intelligence.** (...), **developed the Lisp programming language** (...) **significantly influenced the design of the ALGOL programming language,** (...) and **invented garbage collection.**"
  - in other words... during **Lisp** language design

# Tracing garbage collector



# Tracing garbage collector

- advantages:

# Tracing garbage collector

- advantages:
  - simple to use - no ceremony code (😁)

# Tracing garbage collector

- advantages:
  - simple to use - no ceremony code (😁)
  - "faster" - no overhead during regular work, only occasional pauses

# Tracing garbage collector

- advantages:
  - simple to use - no ceremony code (☺)
  - "faster" - no overhead during regular work, only occasional pauses
  - handles cyclic references out of the box

# Tracing garbage collector

- advantages:
  - simple to use - no ceremony code (☺)
  - "faster" - no overhead during regular work, only occasional pauses
  - handles cyclic references out of the box
- disadvantages:

# Tracing garbage collector

- advantages:
  - simple to use - no ceremony code (☺)
  - "faster" - no overhead during regular work, only occasional pauses
  - handles cyclic references out of the box
- disadvantages:
  - complex - hard to implement, may leak, adds some overhead

# Tracing garbage collector

- advantages:
  - simple to use - no ceremony code (☺)
  - "faster" - no overhead during regular work, only occasional pauses
  - handles cyclic references out of the box
- disadvantages:
  - complex - hard to implement, may leak, adds some overhead
  - non-deterministic deallocation - most (all?) implementations reclaim memory some time **after** an object is no longer needed

# Tracing garbage collector

- advantages:
  - simple to use - no ceremony code (☺)
  - "faster" - no overhead during regular work, only occasional pauses
  - handles cyclic references out of the box
- disadvantages:
  - complex - hard to implement, may leak, adds some overhead
  - non-deterministic deallocation - most (all?) implementations reclaim memory some time **after** an object is no longer needed
  - not completely "pauseless"



THE GARBAGE COLLECTION HAS BEEN CALLED. SOME INTERESTING STATISTICS ARE AS FOLLOWS:

~1960 McCarthy team, MIT symposium

# Tracing garbage collector

Conceptually consists of two simple(™) phases:

# Tracing garbage collector

Conceptually consists of two simple(™) phases:

- **Mark phase** - to discover objects that may be deleted ("trace" them)

# Tracing garbage collector

Conceptually consists of two simple(™) phases:

- **Mark phase** - to discover objects that may be deleted ("trace" them)
- **Collect phase** - to delete discovered objects

# Tracing garbage collector

Conceptually consists of two simple(™) phases:

- **Mark phase** - to discover objects that may be deleted ("trace" them)
- **Collect phase** - to delete discovered objects
  - obviously, the main goal here is not to delete objects themselves, but to reclaim memory after them

# Tracing garbage collector

Conceptually consists of two simple(™) phases:

- **Mark phase** - to discover objects that may be deleted ("trace" them)
- **Collect phase** - to delete discovered objects
  - obviously, the main goal here is not to delete objects themselves, but to reclaim memory after them
  - significant decision is: do we move objects (**Compact**) or not (**Sweep**)

# Tracing garbage collector

Conceptually consists of two simple(™) phases:

- **Mark phase** - to discover objects that may be deleted ("trace" them)
- **Collect phase** - to delete discovered objects
  - obviously, the main goal here is not to delete objects themselves, but to reclaim memory after them
  - significant decision is: do we move objects (**Compact**) or not (**Sweep**)
- both phases may be implemented "concurrently" (to the program execution and/or even to each other) and/or "incrementally"

# Tracing garbage collector - Mark

We need to know which objects are "used"...





# Tracing garbage collector - Mark & Object graph

In memory:



# Tracing garbage collector - Mark & Object graph

In memory:



Type data:

```
record A(B b, D d);  
record B(int X);  
record C(B b, F f);  
record D(E e);  
record E(G g);  
record F(int X);  
record G(int Z);
```

# Tracing garbage collector - Mark & Object graph

In memory:



Type data:

```
record A(B b, D d);  
record B(int X);  
record C(B b, F f);  
record D(E e);  
record E(G g);  
record F(int X);  
record G(int Z);
```

Current "state":

```
var a = new A(..., ...);  
var d = new D(...);  
...we are here...
```

# Tracing garbage collector - Mark & Object graph

In memory:




Type data:

```
record A(B b, D d);  
record B(int X);  
record C(B b, F f);  
record D(E e);  
record E(G g);  
record F(int X);  
record G(int Z);
```


Current "state":

```
var a = new A(..., ...);  
var d = new D(...);  
...we are here...
```


Object graph:



# Object graph traversal




# Object graph traversal




To visit:

# Object graph traversal



To visit: A, D


# Object graph traversal



To visit: A, D, D, B




# Object graph traversal




To visit: ~~A~~, D, D, B

# Object graph traversal




To visit: ~~A~~, D, D, B, E

# Object graph traversal




To visit: ~~A~~, ~~D~~,  $\emptyset$ , B, E

# Object graph traversal




To visit: ~~A~~, ~~D~~,  $\emptyset$ , B, E

# Object graph traversal




To visit: ~~A~~, ~~D~~,  $\emptyset$ , ~~B~~, E, G

# Object graph traversal




To visit: ~~A~~, ~~D~~,  $\emptyset$ , ~~B~~, ~~F~~, G

# Object graph traversal



To visit: ~~A~~, ~~D~~,  $\emptyset$ , ~~B~~, ~~F~~, G


# Object graph traversal




To visit: ~~A~~, ~~D~~,  $\emptyset$ , ~~B~~, ~~F~~,  $\emptyset$




# Object graph traversal



# Object graph traversal




# Object graph traversal




- we have just discovered **reachability** of the objects (from at least one root) by *marking* those reachable.

# Object graph traversal



- we have just discovered **reachability** of the objects (from at least one root) by *marking* those reachable.
- **Reachability** is the closest we can get to true "usability" - we don't know the future.

# Object graph traversal



- we have just discovered **reachability** of the objects (from at least one root) by *marking* those reachable.
- **Reachability** is the closest we can get to true "usability" - we don't know the future.
- objects **C** and **F** may now be deleted by the next phase


# Object graph traversal

Possible roots:

- stack
- CPU registers
- static/thread-local static data
- finalization queue
- inter-generational references ("cards", "card tables")  
(we will return to that...)
- ...


# Collect - Sweep/Compact

## Sweep



# Collect - Sweep/Compact

## Sweep




- unreachable objects are treated as **free space**



# Collect - Sweep/Compact


## Sweep



- unreachable objects are treated as **free space**
  - most will be used as space for new objects

# Collect - Sweep/Compact


## Sweep



- unreachable objects are treated as **free space**
  - most will be used as space for new objects
  - some will become unusable **fragmentation**

# Collect - Sweep/Compact


## Sweep



- unreachable objects are treated as **free space**
  - most will be used as space for new objects
  - some will become unusable
- **fragmentation**
- typically we **DON'T reduce** memory usage - free space is still assigned to the process (☹)

# Collect - Sweep/Compact


## Sweep



- unreachable objects are treated as **free space**
  - most will be used as space for new objects
  - some will become unusable
- **fragmentation**
- typically we **DON'T reduce** memory usage - free space is still assigned to the process (☹)
- fast and simple operation (👍)


# Collect - Sweep/Compact

## Sweep




- unreachable objects are treated as **free space**
  - most will be used as space for new objects
  - some will become unusable
- **fragmentation**
- typically we **DON'T reduce** memory usage - free space is still assigned to the process (☹)
- fast and simple operation (👍)

## Compact



# Collect - Sweep/Compact

## Sweep



- unreachable objects are treated as **free space**
  - most will be used as space for new objects
  - some will become unusable
- **fragmentation**
- typically we **DON'T reduce** memory usage - free space is still assigned to the process (☹)
- fast and simple operation (👍)


## Compact



- move objects so gaps will disappear

# Collect - Sweep/Compact

## Sweep



- unreachable objects are treated as **free space**
  - most will be used as space for new objects
  - some will become unusable
- **fragmentation**
- typically we **DON'T reduce** memory usage - free space is still assigned to the process (🗨)
- fast and simple operation (👍)


## Compact



- move objects so gaps will disappear
- new space for objects is now at the end

# Collect - Sweep/Compact

## Sweep



- unreachable objects are treated as **free space**
  - most will be used as space for new objects
  - some will become unusable
- **fragmentation**
- typically we **DON'T reduce** memory usage - free space is still assigned to the process (🗑️)
- fast and simple operation (👍)

## Compact




- move objects so gaps will disappear
- new space for objects is now at the end
- but... typically **we MAY reduce** memory usage - return this space to the operating system (👍)



# Collect - Sweep/Compact

## Sweep



- unreachable objects are treated as **free space**
  - most will be used as space for new objects
  - some will become unusable
- **fragmentation**
- typically we **DON'T reduce** memory usage - free space is still assigned to the process (☹)
- fast and simple operation (👍)


## Compact



- move objects so gaps will disappear
- new space for objects is now at the end
- but... typically **we MAY reduce** memory usage - return this space to the operating system (👍)
- pretty complex operation with the memory copying overhead (☹)


# Collect - Sweep/Compact

## Sweep



- unreachable objects are treated as **free space**
  - most will be used as space for new objects
  - some will become unusable
- **fragmentation**
- typically we **DON'T reduce** memory usage - free space is still assigned to the process (☹️)
- fast and simple operation (👍)

## Compact



- move objects so gaps will disappear
- new space for objects is now at the end
- but... typically **we MAY reduce** memory usage - return this space to the operating system (👍)
- pretty complex operation with the memory copying overhead (☹️)

**Question:** how to make a decision whether to *sweep* or to *compact*...?! 🤔 We will return to that.

# Automatic memory management

Two biggest lies:

# Automatic memory management

Two biggest lies:

- there is **no** memory leak

# Automatic memory management

Two biggest lies:

- there is **no** memory leak
- there is **no** memory leak unless we use unmanaged resources

# Materials

- <https://play.rust-lang.org>
- [Memory Management Reference](#)
- [The Garbage Collection Handbook: The Art of Automatic Memory Management](#) book (\$)
- [Fundamentals of garbage collection](#)